

Acceleration and variability in the growth of children's early vocabulary

Michael C. Frank

Stanford University


2026

Vocabulary growth

Theoretical: Growth mechanisms are diagnostic of how children's learning works — and how it changes with development.

Practical: Variation in vocabulary growth is a key precursor of academic outcomes.

Methodological: Vocabulary count has a property nearly no other psychological metric has: a true zero.




The data

MacArthur–Bates Communicative Development Inventory (CDI) is a parent-report checklist of ~ 680 words. Parents check which words their child *produces* or *comprehends*.

Wordbank: open repository of CDI administrations. ~ 30 languages, $\sim 10^5$ children. We use both:

- **Cross-sectional:** one snapshot per child
- **Longitudinal:** ≥ 2 admins per child, lets us track per-child growth

Large N , parent-report calibrated to direct testing, cross-linguistically validated.




One child = one point. Vocabulary $\times$ age in English (American) WS.

Two signatures of vocabulary growth

Signature 1: acceleration. Each quantile curve is concave up: doubling age more than doubles vocabulary. The growth rate accelerates. **Signature 2: variability.** The quantile lines fan out: between-child differences in vocabulary grow with age, not shrink.

These signatures are remarkably robust: replicate across languages, samples, and forms. Any model of early word learning must explain both.




The dominant theoretical view: pure accumulation

Each child has input rate r_i ; word j has difficulty δ_j . A word is acquired when enough tokens accumulate.

Textbook lineage:

- McMurray (2007) — pure rate $\times$ time
- Mitchell & McMurray (2009) — + leverage feedback
- Hidaka (2013) — cumulative-Weibull; rate change
- Mollica & Piantadosi (2017) — per-word Gamma
- Kachergis et al. (2021) — IRT Rasch + α_i

All share: between-child variation lives only in baseline (r_i or α_i); per-child scaling exponent fixed.



Each child hears tokens at rate $r_i \cdot p_j$; word j is acquired when its bucket fills past δ_j .


A puzzle: input matters less than supposed

The accumulator view predicts variation in growth comes from variation in input quantity. Coffey & Snedeker (2026) meta-analyzed 71 studies, 4,760 children:

R^2 of input on outcome = 0.04–0.07.

Caregiver speech quantity explains 4–7% of variance in language outcome — a real but small effect.

Hypothesis: Variability comes from *how children accelerate*, not how much input they hear.



Between-child $SD(\theta)$ grows from ~ 1.7 to ~ 2.3 logits between 16 and 30 mo. If input explains $\sim 5\%$, the rest comes from somewhere.

This work: modeling individual variability in acceleration

We build a Bayesian model to estimate three things directly:

1. **Acceleration.** Is population vocabulary growth super-linear in cumulative input?
2. **Variability.** Do children differ in *how* and *when* their growth accelerates, not just in baseline ability?
3. **Input share.** How much of the between-child variance is attributable to input rate, vs. to internal factors like efficiency?

Does the same machinery applied to large language models yield the same answers? **Spoiler:** no — LLMs show a categorical disanalogy with children's learning. (Later in the talk.

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis


The accumulator hypothesis

The shared idea. Each child accumulates word-evidence at her own input rate r_i ; word j is acquired once its difficulty δ_j is crossed.

Ability scales as a power of input. If kids vary only in *rate*, then ability grows as $(\text{tokens})^1$ — the pure *unit accumulator*. Between-child differences come only from differences in r_i .

What our model will check.

- Is the growth exponent really 1, or steeper?
- Are children all alike except for input rate, or do they also differ in efficiency, in growth rate, and in when growth starts?



Bucket-fill metaphor: each child hears tokens at rate $r_i \cdot p_j$; word j is acquired when its bucket fills past difficulty δ_j (Kachergis et al. 2021).

Prior accumulator models: structural lineage

We'll translate each into our notation later, after the model build

McMurray (2007). Pure accumulator. Each word w has acquisition threshold τ_w from population distribution $g(\tau)$. Acceleration requires $g'(\tau) > 0$.

Mitchell & McMurray (2009). McMurray + leverage feedback $a(t) = t + cL(t)$. *Shifts* growth without creating acceleration.

Hidaka (2013). Closed-form AoA distributions: cumulative (Gamma), rate-change (Weibull, exponent D on time), or both.

Mollica & Piantadosi (2017). Per-word Gamma waiting time: $\text{AoA} \sim \Gamma(k_w, \lambda_w)$.

Kachergis et al. (2021). 1-PL Rasch IRT with per-child θ_i and linear age covariate.

Each uses different notation. After we introduce our model, we'll re-express each in our coordinates and compare structurally.

The model: Rasch IRT + accumulator dynamics

$$P(\text{produces}_{ij}) = \sigma(\eta_{ij}) \quad \text{with} \quad \eta_{ij} = \theta_i - \delta_j$$

$$\eta_{ij} = \begin{array}{c} \theta_i \\ \uparrow \\ \text{\textit{kid ability}} \\ \text{(grows with experience)} \end{array} - \begin{array}{c} \delta_j \\ \uparrow \\ \text{\textit{word difficulty}} \\ \text{(per-word, free)} \end{array}$$

This is a standard Rasch IRT model. The substantive content is in how θ_i depends on accumulated experience. We build that up in three steps.

Step 1: pure accumulator

$$\theta_i = \log r_i + \log t \quad \Longleftrightarrow \quad \exp(\theta_i) = r_i \cdot t$$

$$\theta_i = \log r_i + \log t$$



per-child input rate *time = age*
(words per month) *(months since onset)*

Ability = rate × time. The only thing that changes across development is more accumulated input. The only source of variation between children is variation in the input rate r_i .

Step 2: + per-child efficiency α_i

$$\theta_i = \log r_i + \log \alpha_i + \log t \quad \Longleftrightarrow \quad \exp(\theta_i) = \alpha_i \cdot r_i \cdot t$$

$$\theta_i = \log r_i + \log \alpha_i + \log t$$



input rate *per-child efficiency (NEW)* *time*

Each child has a multiplicative efficiency α_i over and above their input rate. **The question we'll ask of the fit:** how much of the between-child variance comes from efficiency vs. input rate?

We summarize this share as $\pi_\alpha = \sigma_\alpha^2 / (\sigma_\alpha^2 + \sigma_r^2)$.

Step 3: + population acceleration κ_{pop}

$$\theta_i = \log r_i + \log \alpha_i + \kappa_{\text{pop}} \log t \quad \Longleftrightarrow \quad \exp(\theta_i) = \alpha_i \cdot r_i \cdot t^{\kappa_{\text{pop}}}$$

$$\theta_i = \log r_i + \log \alpha_i + \kappa_{\text{pop}} \cdot \log t$$



input rate *efficiency* **population
scaling
exponent
(NEW; shared)** *time*


The population accelerates: ability scales as $t^{\kappa_{\text{pop}}}$ super-linearly. All children still share one κ_{pop} .
What we'll ask: is κ_{pop} really above 1 — i.e., does ability grow faster than a unit accumulator predicts? This is Hidaka's rate-change model.

Step 4: + per-child acceleration ζ_i

Children differ not just in baseline ability, but in *how they accelerate*

$$\theta_i = \log r_i + \log \alpha_i + (\kappa_{\text{pop}} + \zeta_i) \log t$$

$$\theta_i = \log r_i + \log \alpha_i + (\kappa_{\text{pop}} + \zeta_i) \cdot \log t$$



efficiency *pop. scaling* **per-child
acceleration
deviation
(NEW)**


Each child has her own growth exponent $\kappa_i = \kappa_{\text{pop}} + \zeta_i$. The new piece: between-child variation isn't just in baseline ability, but in *how trajectories accelerate*. Every prior accumulator model forces every child to share one growth rate. **What we'll ask:** do children really differ in how fast their vocabulary accelerates?

Step 5: + per-child trajectory phase s_i

The second piece of *individual variability in acceleration*: when children start

$$\theta_i = \log r_i + \log \alpha_i + (\kappa_{\text{pop}} + \zeta_i) \log(t - s - s_i)$$

$$\theta_i = \dots + (\kappa_{\text{pop}} + \zeta_i) \cdot \log \left(\begin{array}{c} t \quad - s - \quad s_i \end{array} \right)$$


calendar age **per-child
trajectory
phase (NEW)**

The same kind of individual-variability addition, but on the *time* dimension: children differ in *when* their developmental clock effectively starts. s_i can be positive (later start) or negative (developmentally ahead). **What we'll ask:** does the timing of growth onset really vary across children, alongside how fast they accelerate?

What every prior model pins; what we free

Re-expressing each prior model in our common notation,

$$\eta_{ij}(t) = \log r_i + \log \alpha_i + \kappa_i \log(t - s - s_i) - \delta_j$$

Model	efficiency	pop. growth	per-child	per-child	Per-word
	σ_α	κ_{pop}	growth σ_κ	phase σ_s	δ_j ?
McMurray (2007)	0	1	0	0	free, $\delta_j = \log \tau_j$
Mitchell & McMurray (2009)	0	1*	0	0	free, plus leverage c
Hidaka (2013) cumulative	0	1	0	0	free, $\delta_j = \log(N_j/\delta)$
Hidaka (2013) rate-change	0	$D+1$	0	0	free
Mollica & Piantadosi (2017)	0	1	0	0**	free, $\delta_j = \log(k_j/\lambda_j)$
Kachergis et al. (2021)	free	free	0	0	free
Best-fitting model (this work)	free	free	free	free	free

What's structurally new. Every prior model pins σ_κ and σ_s to zero (same growth rate and onset for all kids). Our model frees both; we find $\sigma_\kappa \approx 3.5$ and $P(\kappa_i > 1) > 0.997$ — the load-bearing additions that let the model fit acceleration and variability *jointly*.

* Mitchell & McMurray's leverage isn't linear in $\log t$. ** Mollica's Gamma encodes a per-word offset, but shared

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis

Where does σ_r come from? The input-rate prior

Why we need an external pin: the variance partition is not data-identified

The kid-level intercept $\xi_i = \log r_i + \log \alpha_i$ has total variance σ_ξ^2 that the data *does* identify. But the split between input rate (σ_r^2) and efficiency (σ_α^2) is not separately identified within the model:

$$\sigma_\xi^2 = \sigma_r^2 + \sigma_\alpha^2 \quad \Rightarrow \quad \pi_\alpha = 1 - \frac{\sigma_r^2}{\sigma_\xi^2}$$

We resolve this by pinning σ_r externally — from data quantifying actual variation in child-directed input rates across families. The model then partitions the data-identified σ_ξ^2 into input and efficiency components.

The pin matters for π_α .

- Bigger $\sigma_r \Rightarrow$ more variance assigned to input $\Rightarrow$ lower π_α .
- Pinning σ_r *too low* would overstate efficiency's share; *too high* would mask it.

Strategy. Compile the empirical literature on adult-token rates per child, derive a defensible σ_r , and report a full sensitivity sweep so the headline doesn't hinge on the central pin.

Empirical estimates of input-rate variation

From observational studies of adult speech to young children

Group	n	Mean (tok/hr)	± 1 SD range (tok/hr)	Tokens/month mean (range)	σ_r (log)
Western, CDS-mother (Sperry sites, HR, WF)	17	$\sim 1,100$	620–2,000	400K (230K–730K)	0.57
Western, all-adult tokens	9	$\sim 2,200$	1,450–3,400	800K (530K–1.24M)	0.43
Sperry within-pool, CDS-any-adult	42	$\sim 1,500$	880–2,620	547K (320K–956K)	0.53
All groups (incl. Tseltal, Yélî, Pirahã)	27	—	~ 300 –5,000	(~ 110 K–1.8M)	1.24


Central pin: $\sigma_r = 0.534$ from the Sperry within-pool. A typical Western child hears ~ 547 K adult tokens/month (± 1 SD: ~ 320 K–960K, a $2.5\times$ spread). The Wordbank CDI:WS sample is overwhelmingly WEIRD, so the cross-cultural 1.24 is an outer bound, not a likely value.

Tokens/month = tokens/hr $\times$ H, H = 365 hr/mo (12 waking hr/day $\times$ 30.44 days/mo, model convention).

Sources: Sperry et al. 2019, Hart & Risley 1995, Weisleder & Fernald 2013, Bergelson 2018/2019/2023, Bunce et al. 2024, Casillas 2020/2021, Coffey et al. 2024.

Sensitivity: π_α across plausible σ_r

Stress-testing the input-rate pin



$\pi_\alpha(\sigma_r) = 1 - \sigma_r^2/\sigma_\xi^2$, with σ_ξ fixed at its data-identified posterior.

Central pin $\sigma_r = 0.534$: $\pi_\alpha \approx 0.91-0.94$.

Doubled $\sigma_r = 1.2$ (cross-cultural upper bound): π_α still > 0.55 .

Halved $\sigma_r = 0.30$ (below any defensible single-site SD): $\pi_\alpha \approx 0.97$.

Conclusion. The efficiency-dominated decomposition holds across the full range supported by Western input data, and the qualitative answer survives even the cross-cultural upper bound.

Caveat: this is about input *quantity*, not *quality*

What the σ_r analysis does and doesn't address

What we've modeled. Variation in input *quantity* — how many adult tokens per hour each child hears. Under this lens, between-child differences in vocabulary growth are mostly *not* explained by input quantity.

What we haven't modeled. Variation in input *quality* — diversity, contingency, conversational structure, child-directed register. Quality variation is real and may matter for learning.

Why this is still a defensible scope:

- **Quality is hard to measure at scale.** No cross-population benchmark analogous to LENA-style token rates.
- **Quality would have to be enormous to escape the finding.** A per-child age-varying quality multiplier $q_i(t) = (t/a_0)^{\gamma_i}$ is observationally equivalent to a steeper κ_i — so to absorb the observed acceleration, “quality” would need to grow $\sim 3\{, \}000\times$ for the median kid and $\sim 10^6\times$ for the 95th percentile between 12 and 30 mo. (See appendix.)

Bottom line. Quantity is too small and too well-measured to explain the variability; quality would have to be implausibly large and age-varying to substitute.

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis

Data and compute for the longitudinal fits

How the next results were produced

Data.

- **English Wordbank (WS form).** $I = 500$ children, ~ 1500 administrations, $J = 671$ items (full WS list), $N \approx 1.1\text{M}$ observed responses. Stratified-by-median-age sampling.
- **Norwegian Wordbank.** $I = 500$ children, ~ 3000 administrations, $J = 716$ items, $N \approx 2.2\text{M}$ responses. Same stratification.


Model family. The same five-step build run as a sequence of Stan models on each dataset: pure, $+\alpha$, $+\kappa_{\text{pop}}$, $+\zeta_i$, $+s_i$. Each fit warm-starts from the previous one's posterior.

Compute. Two GCP virtual machines (n2d-standard-128, 64 physical AMD EPYC cores each, 128 GB RAM), one per language, running in parallel. Stan with `reduce_sum` threading: 4 chains $\times$ 16 threads/chain saturates the 64 physical cores per machine. $\sim 11\text{h}$ per fit, $\sim 55\text{h}$ per VM for the five-variant family; the two VMs run in parallel, so $\sim 55\text{h}$ wall-clock and $\sim 110\text{h}$ of 128-core VM time total. No GPUs.

Priors and pins. $s = 6$ mo (population trajectory anchor); $\sigma_r = 0.534$ (Western-CDS empirical pin, see previous slides). Posteriors with $\hat{R} < 1.05$ on all parameters of interest.

Each addition matters: 5-panel comparison to data

Variants fit to Wordbank English WS; predicted percentile fans vs. observed




Reading. Each panel adds one mechanism. Pure accumulator collapses kids to one trajectory; σ_α adds vertical spread; κ_{pop} adds curvature; σ_ζ opens the fan with age; σ_s handles the floor at young ages. 25 / 48

Same story in θ -space

The structural commitment each model makes about per-child latent ability

Predicted latent ability across the 5-model build

Each panel: 100 simulated kids (grey), population mean (black), 10th/90th percentiles (blue/red dashed). Dotted: θ where a typical-difficulty word has 50% acquisition probability. Posterior values from current English fits ($s=0.5$); will refresh after $s=6$ refit.



Same five-panel build, but plotting latent ability $\theta_i(t)$ directly instead of integrating to vocabulary. Population mean (bold), 10th/90th percentiles (blue/red), 100 simulated kids (faint). Dotted: $\theta_{50\%}$, the value at which a typical-difficulty word has 50% acquisition probability.

What it shows. The first two panels never cross $\theta_{50\%}$ within the data range — pure accumulators (with or without efficiency variation) are structurally unable to fit children's acquisition. Each subsequent panel adds a structural mechanism the data demands.

Best-fitting model: parameter posteriors

Acceleration and variability both bounded firmly above zero

($s=0.5$ placeholders)

Parameter	Posterior	What it does
σ_α	1.71 [1.55, 1.91]	Between-child efficiency variation.
π_α	0.91 [0.89, 0.93]	Share of intercept variance from efficiency, not input.
κ_{pop}	10.35 [10.21, 10.51]	Population scaling exponent: ability scales as (tokens) $^{\kappa_{\text{pop}}}$.
σ_ζ	3.47 [3.14, 3.84]	Between-child variation in scaling exponent.
σ_s	[refits pending]	Between-child variation in trajectory phase.
$\rho(\xi, \zeta)$	+0.35 [+0.22, +0.47]	Higher-baseline kids have steeper growth.

Variability \[0.2em] $\pi_\alpha = 0.91$:
efficiency variance dominates
input variance.

Acceleration \[0.2em]
 $\kappa_{\text{pop}} \approx 10.4 \gg 1$: scaling is
wildly super-linear, far from
the unit accumulator.

Variation in scaling \[0.2em]
 $\sigma_\kappa > \sigma_\alpha$: kids vary more in scaling
exponent than in baseline.

Cross-language acceleration

Where it lives, not whether it does

Sample	admins/kid	κ_{pop}	σ_{κ}	kid range ($\kappa_{\text{pop}} \pm 2\sigma_{\kappa}$)
English	median 3	10.4 [10.2, 10.5]	3.5 [3.1, 3.8]	~ 3 to 17
Norwegian	median 8	12.4 [12.3, 12.6]	3.7 [3.4, 4.2]	~ 5 to 20

$\kappa_{\text{pop}} \sim 10-12$ in both languages, with
 $P(\kappa_i > 1) > 0.997$.

δ vs. ζ_i **decomposition is identifiability-dependent.** In Norwegian (8 admins/kid), per-child slopes are well-identified — dropping ζ_i hurts fit by 235 ELPD. In English (3 admins/kid), δ pools what Norwegian's data lets fall into ζ_i .

Reading. The acceleration signature is universal. Its internal parameterization is identifiability-dependent on data density.

In any dense longitudinal sample, between-child κ_i spans roughly 5–15.

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis

Three model extensions: testing π_α with richer data

The Wordbank longitudinal fits use *population-level* input rate priors (Sperry/Hart-Risley/Weisleder-Fernald). Two natural follow-ups:

1. **Per-child observed input:** use *measured* per-child input rate instead of a population prior. Does π_α stay ~ 0.9 when input is no longer a pooled estimate? → **BabyView, SEEDLingS.**
2. **Per-child processing efficiency:** link CDI growth to an independent per-child readout of processing speed. Does the σ_α we estimate from CDI reflect “processing efficiency” or something else? → **Peekbank.**

Both extensions reuse the same model machinery with an added channel in the linear predictor; only the data and the per-child latents change.


Replicating with directly observed input

BabyView (head-mounted video) and SEEDLingS (LENA all-day audio)

BabyView, $N = 20$

Per-child observed input vs inferred true rate

Dashed green: $y = x + \text{beta_react}$ ($= 0.28$); dotted: $y = x$. Reactivity inflation = 33% (CrI: -35-173%)




$$\pi_{\alpha} = 0.84, [0.68, 0.93]$$

SEEDLingS, $N = 44$ (comp corrected)

Per-child observed input vs inferred true rate

Dashed green: $y = x + \text{beta_react}$ ($= -0.00$); dotted: $y = x$. Reactivity inflation = -0% (CrI: -0-0%)



$$\pi_{\alpha} = 0.94, [0.89, 0.97]$$

With per-recording observed input, the π_{α} headline holds: input-rate variance is a small share of between-child variance even when measured directly.

Adding a processing-speed readout: Peekbank LWL

$N = 62$ subjects with both CDI and looking-while-listening

Joint model adds a per-child LWL log-RT readout coupled to $\log \alpha_i$ via γ_{rt} .

Posterior.

- $\gamma_{rt} = 0.082$, $[0.046, 0.125]$ — LWL processing speed is real, and bounded above zero.
- $\pi_{\alpha} = 0.90$, $[0.86, 0.94]$ — replicates within the Peekbank subset.
- $\rho(\kappa_i, \text{rtslope}_i) = -0.02$, $[-0.33, +0.27]$ — CDI-side and LWL-side growth rates do *not* share variance.

Two clocks, not one. The acceleration in vocab growth and the maturation of processing speed are independent maturation processes within the same child — different “efficiency clocks” ticking in parallel.

This rules out the simplest unification (“one underlying efficiency axis”); the two readouts each pick up a different developmental process.

π_α across five samples

Three input-observation channels, two languages, parent-report noise corrected

Sample	N	σ_α	π_α	Channel
English longitudinal	200	1.83 [1.65, 2.04]	0.91 [0.90, 0.93]	Pooled input prior
Norwegian longitudinal	200	2.10 [1.90, 2.35]	0.94 [0.93, 0.95]	Pooled input prior
Peekbank (Stanford)	62	1.64 [1.34, 2.03]	0.90 [0.86, 0.94]	Pooled + LWL
BabyView	20	1.13 [0.86, 1.61]	0.84 [0.68, 0.93]	Head-mounted video
SEEDLingS	44	1.39 [1.13, 1.74]	0.94 [0.89, 0.97]	LENA all-day audio

$\pi_\alpha \in [0.84, 0.94]$ **across all five samples.** The specific input channel (pooled vs. measured) and sample size hardly move the picture: between-child variance is dominated by something *other* than input rate.

(Placeholder: numbers from $s=0.5$ fits; $s=6$ refits running. The qualitative pattern is invariant.)

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis

Why look at LLMs?

Same task, same statistical machinery, very different learners

LLMs and children both learn word meanings from input. They differ *enormously* in scale and architecture, but the question of how their learning scales with input is statistically well-posed:

- For an LM and a word w , we have $P(w \text{ produced correctly})$ as a function of training tokens.
- This is a sigmoid in log-tokens with a per-word slope.
- The same statistic exists for human children: $P(w \text{ acquired})$ as a function of cumulative input.

This lets us put kids and LMs on the same axis — not just trade anecdotes about scaling.
Two questions:

1. Do kids and LMs have similar per-word acquisition slopes?
2. Does a model that fits kids' learning say anything about LMs?

Chang & Bergen (2022): per-word sigmoids in LMs

Method we'll re-use for the comparison

Setup. Chang & Bergen trained four LMs (BERT, biLSTM, GPT-2, LSTM) on the same input corpus (BookCorpus + WikiText). At each of ~ 200 training checkpoints, they probed each of ~ 611 CDI words to get $P(\text{word produced correctly} \mid \text{checkpoint})$.

For each (LM, word) they fit a 4-parameter logistic

$$S_w(x) = \text{lower}_w + \frac{\text{upper}_w - \text{lower}_w}{1 + \exp((x - \text{mid}_w)/\text{scale}_w)}$$

where $x = \log_{10}(\text{training steps})$, yielding a per-word **scaling slope**: how fast that word's acquisition probability rises with cumulative training. After unit conversion, this is exactly analogous to our $\kappa_i \log t$ for children.

Result they reported. Across the four LMs, mean per-word slope ≈ 1 logit per nat-log training-step. Striking but ambiguous: is this the LM architecture or the input distribution?

Methods: re-running C&B on CHILDES-trained LMs

Feng et al. (2026) trained GPT-2 on the input distribution kids actually hear

Confound in C&B's setup: BookCorpus + WikiText is adult written text, not child-directed speech. That conflates two things: (a) structural difference in how LMs vs children scale; (b) input distribution mismatch.


Feng et al. (2026) train multiple GPT-2 variants on *CHILDES* — the same corpus that supplies our $\log p_j$ frequencies for the kid models. Same architecture as Chang & Bergen's GPT-2, but input matched to what children actually hear.

Our re-application: for each Feng et al. checkpoint, fit C&B's per-word sigmoid on ~ 611 CDI words. Compute per-word slope in our κ -equivalent units. Compare to kids.

Status: we refit C&B's per-word sigmoids on the Feng et al. CHILDES-trained GPT-2 checkpoints (Sherlock GPU jobs). The CHILDES-trained slopes are what appear in the comparison on the next slide — not the BookCorpus original.

The per-word view: are kids and LMs learning at the same rate?

Per-instance sigmoid slope on the same x-axis (log experience)





Reading. Per-word sigmoid slope on $\log(\text{experience})$: LMs (BERT, biLSTM, GPT-2, LSTM, from Chang & Bergen 2022 + our CHILDES-trained refits) vs. our per-child κ_i . **LMs cluster near $\kappa \approx 1$;** children average $\approx 10\times$ steeper, with much more between instance variance

The aggregate view: scaling laws

Children sit far above LM scaling-law slopes on cumulative tokens

Slopes are the structurally meaningful comparison; y-axis offset is anchored for visual clarity. Kid D range comes from per-kid r_i .



SM2 fit long_slopes : median $1+\delta = 10.39$ [10.24, 10.55]; $\sigma_\zeta = 3.48$ [3.15, 3.87]; $\mu_{u_r} = 7.34$, $\sigma_{u_r} = 0.53$, $a_u = 19$. Chinchilla and Kaplan exponents from Hoffmann et al. 2022 / Kaplan et al. 2020.

Comparing two estimates of *ability gained per token of experience*: LM scaling laws (Kaplan, Chinchilla compute-optimal, Chinchilla data-bound) all have exponents $\beta \in [0.10, 0.28]$. Children's $\kappa_{\text{pop}} \approx 10$ is two orders of magnitude steeper — a different scaling regime, not a quantitative degree

Outline

1. Building the model: individual variability in acceleration
2. Input quantity: how much does it really vary?
3. Results: English and Norwegian longitudinal fits
4. Extensions: directly observed input and processing
5. Connecting to LLMs
6. Synthesis

Technical summary


What the modeling pinned down

- **Acceleration is real and population-level.** $\{(placeholder\ for\ s=6\ refits)\}$ $\kappa_{pop} \approx 10$, confidently above the unit-accumulator boundary $\kappa_{pop} = 1$. The first prior model to free κ_{pop} (Hidaka rate-change) was structurally right.
- **Individual variability in acceleration is real too.** $\sigma_{\zeta} > 0$ and $\sigma_s > 0$: children differ not just in baseline ability, but in *how their trajectories accelerate* and *when*. No prior accumulator model frees these.
- **Input quantity is a small share of between-child variance.** $\pi_{\alpha} \approx 0.91$ across 5 samples. Aligns with the meta-analytic $R^2 = 0.04\text{--}0.07$ of Coffey & Snedeker (2026).
- **LLMs are structurally different.** Same per-word sigmoid statistic gives $\kappa \sim 1$ for LMs vs. ~ 10 for children — a $10\times$ gap with near-zero between-instance variance in LMs.

github.com/langcog/standard-model-2


Two signatures, twin constraints on theory

Calling back to the start



Acceleration. $\kappa_{\text{pop}} \approx 10$ — above any pure accumulator. *Where does the super-linearity come from?*


Twin constraints on any theory: **(a)** population-level super-linear growth (maturation? learning-to-learn?); **(b)** individual differences *not* in input quantity (processing? interactional quality?). *Pure accumulators satisfy neither.*



Variability. $\pi_{\alpha} \approx 0.91$: input explains $\sim 10\%$ of between-child variance. *Where does the other 90% live?*

Appendix: per-child trajectories from the best fit

Acceleration and variability both visible at the level of individual kids




120 kids sampled from the best-fitting model's posterior. **Left:** latent ability θ (logit). **Right:** predicted vocabulary. Concave-up shape = super-linear signature; fan width = between-child variability

Appendix: σ_r sensitivity from earlier cross-sectional fits

Confirms the variance-partition story doesn't hinge on the central pin

Quality-multiplier interpretation: same model, different label



$\delta + \zeta_{\text{token}} \sim N(9.35, 3.47^2)$. For 'token quality grows with age' to absorb M_best's acceleration, median kid needs $q(30)/q(12) \sim 3377x$; ± 1 SD spread covers $40x$ to $981046x$.

Reviewer reflex. “Maybe kids’ tokens differ in *quality*, and that’s what looks like acceleration.”

Structural answer. Per-child age-varying quality $q_i(t) = (t/a_0)^{\gamma_i}$ on input is observationally equivalent to $\gamma_i = \kappa_i - 1$.

Implication. Quality cannot escape acceleration — it *relabels* it. For the median kid, “quality” would have to grow $\sim 3,000\times$ from age 12 to 30 mo. For the 95th percentile kid, $\sim 10^6\times$. Implausible as a substantive quality story.

Appendix: McMurray (2007) in our notation

Pure accumulator: all four parameters pinned at the unit-rate boundary

Original. Each word w has threshold $\tau_w \sim g(\tau)$; deterministic step at $t \geq \tau_w$. Vocabulary $L(t) = G(t)$ (CDF of g). Acceleration requires $g'(\tau) > 0$ over the productive window.

In our notation:

$$\eta_{ij}(t) = \log t - \log \tau_j = \kappa \log t - \delta_j \quad \text{with} \quad \kappa \equiv 1, \quad \delta_j = \log \tau_j.$$

Parameter	Setting in McMurray
$\log r_i$	$\equiv 0$ (no per-child rate variation)
$\log \alpha_i$	$\equiv 0$ (no efficiency variation)
κ_i	$\equiv \mathbf{1}$ (unit accumulator)
δ_j	free, drawn from g

Appendix: Hidaka (2013) in our notation

Rate-change form: D is exactly $\kappa - 1$

Original. Three families: cumulative (N -accumulator at constant rate, $\text{AoA} \sim \Gamma$), rate-change (1-accumulator at rate $\delta \cdot t^D$, $\text{AoA} \sim \text{Weibull}$), and both. **Rate-change:** ability $\propto t^{D+1}$; $D > 0$ produces super-linear scaling.

In our notation (rate-change form):

$$\eta_{ij}(t) = (D + 1) \log t - \delta_j \quad \text{with} \quad \kappa_i \equiv D + 1.$$

Parameter	Setting in Hidaka rate-change
$\log r_i$	$\equiv \log \delta_{\text{const}}$ (shared population rate)
$\log \alpha_i$	$\equiv 0$ (no efficiency variation)
κ_i	$\equiv \mathbf{D + 1}$ (shared, no per-child variation)
δ_j	free

Hidaka's $D = \kappa_{\text{pop}} - 1$ in our notation. His $D > 0$ is our $\kappa_{\text{pop}} > 1$.

Appendix: Mollica & Piantadosi (2017) in our notation

Per-word Gamma waiting time, no per-child structure

Original. Per-word: k_w effective learning instances arrive as Poisson at rate λ_w ; word acquired after k_w events. $AoA_w \sim \Gamma(k_w, \lambda_w)$. For large k_w , Gamma CDF approximates Gaussian (mean k_w/λ_w) and then a ~ 1.7 -scale logistic.

In our notation:

$$\eta_{ij}(t) = \log t - \delta_j \quad \text{with} \quad \kappa_i \equiv 1, \quad \delta_j \approx \log(k_j/\lambda_j).$$

Parameter	Setting in Mollica & Piantadosi
$\log r_i$	$\equiv \log \lambda$ (shared population rate)
$\log \alpha_i$	$\equiv 0$ (no per-child variation)
κ_i	$\equiv \mathbf{1}$ (unit-rate accumulator)
δ_j	$\approx \log(k_j/\lambda_j)$, free per-word

Appendix: Kachergis et al. (2021) in our notation

First prior model to free σ_α ; σ_κ still pinned

Original. 1-PL Rasch IRT with per-child latent ability:

$\text{logit } P(y_{ij} = 1) = \theta_i + d_j + \beta_a a_i + \beta_t \log(p_j H a_i)$. Per-child θ_i free; d_j per-word free; shared age slope β_a across all kids.

In our notation:

$$\eta_{ij} = (\log r_i + \log \alpha_i) + \kappa \log a_i - \delta_j \quad \text{with} \quad \kappa \equiv \text{shared population slope.}$$

Parameter	Setting in Kachergis 2021
$\log r_i$	free, no per-child variance specified
$\log \alpha_i$	free (their θ_i absorbs $\log r + \log \alpha$)
κ_i	$\equiv$ shared population slope (no per-child variance)
δ_j	free (their d_j)

First prior model to free σ_α . But σ_κ is still zero — every kid has the same scaling slope. 48 / 48